

```
In [1]: ► import pandas as pd
import numpy as np

import tensorflow as tf

import keras

import matplotlib.pyplot as plt
import seaborn as sns
```

Using TensorFlow backend.

```
In [2]: ► ibm_stock_data = pd.read_csv('../data/ibm_stock_prices.csv')
```

```
In [3]: ► ibm_stock_data.shape
```

Out[3]: (14059, 7)

```
In [4]: ► ibm_stock_data.head()
```

Out[4]:

	Date	Open	High	Low	Close	Volume	OpenInt
0	1962-01-02	6.4130	6.4130	6.3378	6.3378	467056	0
1	1962-01-03	6.3378	6.3963	6.3378	6.3963	350294	0
2	1962-01-04	6.3963	6.3963	6.3295	6.3295	314365	0
3	1962-01-05	6.3211	6.3211	6.1958	6.2041	440112	0
4	1962-01-08	6.2041	6.2041	6.0373	6.0870	655676	0

```
In [5]: ► ibm_stock_data.drop('OpenInt', axis=1, inplace=True)
```

```
In [6]: ▶ ibm_stock_data.isnull().any()
```

```
Out[6]: Date      False
        Open      False
        High      False
        Low       False
        Close     False
        Volume    False
        dtype: bool
```

Convert data to number

```
In [7]: ▶ ibm_stock_data.dtypes
```

```
Out[7]: Date      object
        Open     float64
        High     float64
        Low      float64
        Close    float64
        Volume   int64
        dtype: object
```

```
In [8]: ▶ ibm_stock_data['Date'] = ibm_stock_data['Date'].astype('datetime64')
```

```
In [9]: ▶ ibm_stock_data.head()
```

```
Out[9]:
```

	Date	Open	High	Low	Close	Volume
0	1962-01-02	6.4130	6.4130	6.3378	6.3378	467056
1	1962-01-03	6.3378	6.3963	6.3378	6.3963	350294
2	1962-01-04	6.3963	6.3963	6.3295	6.3295	314365
3	1962-01-05	6.3211	6.3211	6.1958	6.2041	440112
4	1962-01-08	6.2041	6.2041	6.0373	6.0870	655676

```
In [10]: ▶ ibm_stock_data.dtypes
```

```
Out[10]: Date      datetime64[ns]
Open      float64
High      float64
Low       float64
Close     float64
Volume    int64
dtype: object
```

```
In [11]: ▶ min_date = ibm_stock_data['Date'].min()
print("min_date : ", min_date)
```

```
min_date : 1962-01-02 00:00:00
```

```
In [12]: ▶ max_date = ibm_stock_data['Date'].max()
print("max_date : ", max_date)
```

```
max_date : 2017-11-10 00:00:00
```

```
In [13]: ▶ def calculate_days_since_min_date(row_date):
return row_date - min_date
```

```
In [14]: ▶ ibm_stock_data['numeric_date'] = ibm_stock_data['Date'].apply(lambda x: calculate_days_since_min_date(x))
```

```
In [15]: ▶ ibm_stock_data.head()
```

```
Out[15]:
```

	Date	Open	High	Low	Close	Volume	numeric_date
0	1962-01-02	6.4130	6.4130	6.3378	6.3378	467056	0 days
1	1962-01-03	6.3378	6.3963	6.3378	6.3963	350294	1 days
2	1962-01-04	6.3963	6.3963	6.3295	6.3295	314365	2 days
3	1962-01-05	6.3211	6.3211	6.1958	6.2041	440112	3 days
4	1962-01-08	6.2041	6.2041	6.0373	6.0870	655676	6 days

```
In [16]: ▶ ibm_stock_data.dtypes
```

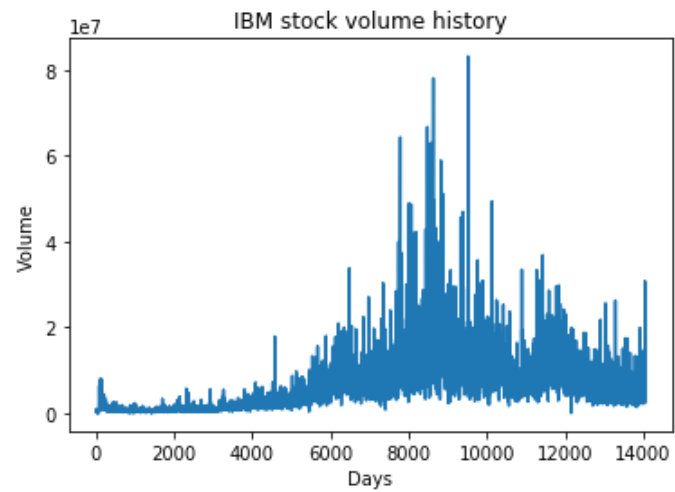
```
Out[16]: Date          datetime64[ns]  
Open              float64  
High              float64  
Low               float64  
Close             float64  
Volume            int64  
numeric_date      timedelta64[ns]  
dtype: object
```

```
In [17]: ▶ ibm_stock_data['numeric_date'] = ibm_stock_data['numeric_date'].astype('timedelta64[D]')
```

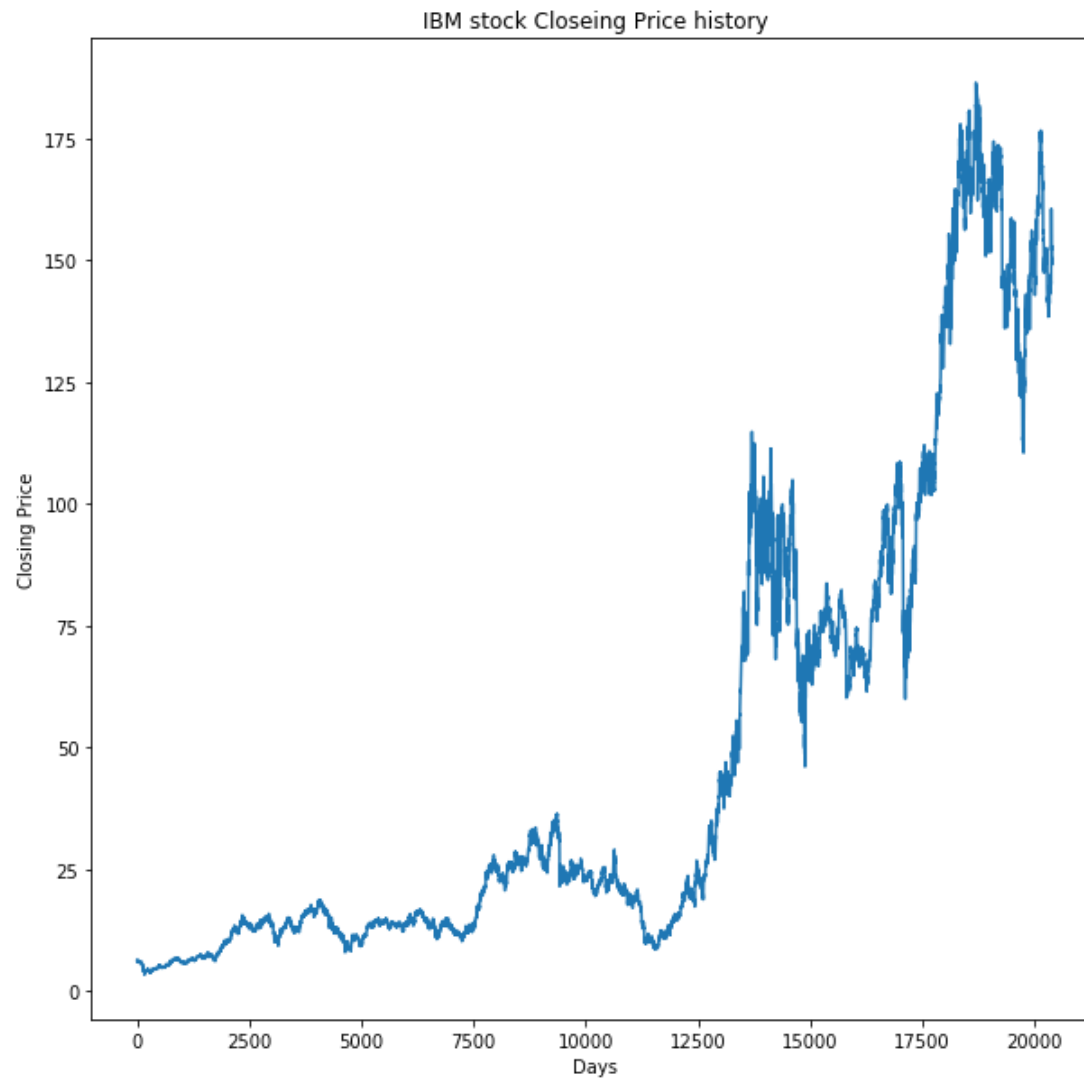
```
In [18]: ▶ ibm_stock_data.dtypes
```

```
Out[18]: Date          datetime64[ns]  
Open              float64  
High              float64  
Low               float64  
Close             float64  
Volume            int64  
numeric_date      float64  
dtype: object
```

```
In [19]: ▶ plt.figure()  
plt.plot(ibm_stock_data["Volume"])  
plt.title('IBM stock volume history')  
plt.ylabel('Volume')  
plt.xlabel('Days')  
plt.show()
```



```
In [20]: ▶ plt.figure(figsize=(10,10))
plt.plot(ibm_stock_data["numeric_date"], ibm_stock_data["Close"])
plt.title('IBM stock Closeing Price history')
plt.ylabel('Closing Price')
plt.xlabel('Days')
plt.show()
```



```
In [21]: ▶ ibm_stock_data.drop('Date', axis=1, inplace=True)
```

```
In [22]: ▶ ibm_stock_data['numeric_date'].min()
```

```
Out[22]: 0.0
```

```
In [23]: ▶ ibm_stock_data['numeric_date'].max()
```

```
Out[23]: 20401.0
```

```
In [24]: ▶ TIME_STEPS = 60
```

```
In [25]: ▶ stock_train_set = ibm_stock_data[ibm_stock_data.numeric_date <= 19000.0]  
stock_test_set = ibm_stock_data[ibm_stock_data.numeric_date > 19000.0]
```

```
In [26]: ▶ stock_train_set.shape
```

```
Out[26]: (13091, 6)
```

```
In [27]: ▶ stock_train_set = stock_train_set.iloc[0:13020]  
stock_test_set = stock_test_set.iloc[0:900]
```

```
In [28]: ▶ stock_train_set.shape
```

```
Out[28]: (13020, 6)
```

```
In [29]: ▶ #stock_test_set[:100]
```

```
In [30]: ▶ stock_test_set.shape
```

```
Out[30]: (900, 6)
```

```
In [31]: ▶ X_train_orig = stock_train_set[['Open', 'High', 'Low', 'Volume', 'numeric_date']].values  
#X_train = stock_train_set[['Open', 'High', 'Low', 'Volume']].values  
#X_train = stock_train_set[['Open', 'High', 'Low', 'Volume', 'numeric_date']].values  
y_train = stock_train_set['Close'].values
```

```
In [32]: ▶ X_test_orig = stock_test_set[['Open', 'High', 'Low', 'Volume', 'numeric_date']].values
#X_test = stock_test_set[['Open', 'High', 'Low', 'Volume']].values
#X_test = stock_test_set[['Open', 'High', 'Low', 'Volume', 'numeric_date']].values
y_test = stock_test_set['Close'].values
```

```
In [33]: ▶ from sklearn.preprocessing import MinMaxScaler

sc = MinMaxScaler()
X_train = sc.fit_transform(X_train_orig)
X_test = sc.transform(X_test_orig)
```

```
In [34]: ▶ X_test.shape
```

```
Out[34]: (900, 5)
```

```
In [35]: ▶ a = [0,1,2,3,4,5,6,7,8,9]
print(a[:-2])
print(a[2:])
print('(0,1:2), (1,2:3), (2,3:4), (3,4:5), (4,5:6), (5,6:7),(6,7:8), (7,8:9)')
```

```
[0, 1, 2, 3, 4, 5, 6, 7]
[2, 3, 4, 5, 6, 7, 8, 9]
(0,1:2), (1,2:3), (2,3:4), (3,4:5), (4,5:6), (5,6:7),(6,7:8), (7,8:9)
```



```

In [36]: ▶ # Each TIME_STEP is a row 'Open', 'High', 'Low', 'Volume', 'numeric_date', we are considering 4 of these to predict 5th.
# Hence number of TIME_STEPS = 4

def build_time_series_data(x, y):

    #iters = x.shape[0]//TIME_STEPS

    row_dim = TIME_STEPS
    col_dim = x.shape[1]
    third_dim = x.shape[0] - TIME_STEPS

    print(third_dim, row_dim, col_dim)

    x_time_series_data = np.zeros((third_dim, row_dim, col_dim))
    y_time_series_data = np.zeros((third_dim))
    num_date_corresponding_to_y = np.zeros((third_dim))

    for i in range(third_dim):
        x_time_series_data[i] = x[i:TIME_STEPS+i]
        y_time_series_data[i] = y[TIME_STEPS+i]

        num_date_corresponding_to_y[i] = sc.inverse_transform(x)[TIME_STEPS+i][4] # as 5th columns in X is the "numeric_date"

    return x_time_series_data, y_time_series_data, num_date_corresponding_to_y

```

```

In [37]: ▶ x_train_time_series_data, y_train_time_series_data, train_num_date_corresponding_to_y = build_time_series_data(X_train, y_train)
x_test_time_series_data, y_test_time_series_data, test_num_date_corresponding_to_y = build_time_series_data(X_test, y_test)

```

```

12960 60 5
840 60 5

```

```

In [38]: ▶ y_test_time_series_data[y_test_time_series_data > 0].shape

```

```

Out[38]: (840,)

```

```

In [39]: ▶ y_train_time_series_data[y_train_time_series_data != 0].shape

```

```

Out[39]: (12960,)

```

```
In [40]: ▶ from keras import Sequential
from keras.layers import GRU
from keras.layers import LSTM
from keras.layers import Dense
from keras.layers.core import Dropout
from keras.layers.core import Flatten
from keras.layers import GRU
from keras import optimizers

from keras import regularizers
```

```
In [41]: ▶ BATCH_SIZE = 60

t_row_dim = TIME_STEPS
t_col_dim = X_train.shape[1]
t_third_dim = X_train.shape[0] - TIME_STEPS
```

```
In [42]: ► rnn_model = Sequential()
rnn_model.add(GRU(3, return_sequences=True, input_shape=(t_row_dim, t_col_dim) ))
#rnn_model.add(Dropout(0.7))

rnn_model.add(GRU(3, return_sequences=False))
#rnn_model.add(Dropout(0.7))

#rnn_model.add(Flatten())
#rnn_model.add(Dense(25))
#rnn_model.add(Dropout(0.6))

rnn_model.add(Dense(5))
rnn_model.add(Dropout(0.3))

rnn_model.add(Dense(1))

optimizer = optimizers.Adam(lr=0.7)
rnn_model.compile(optimizer = 'adam', loss = 'mean_squared_error')
```

WARNING:tensorflow:From C:\Users\thisi\AppData\Local\Continuum\anaconda3\lib\site-packages\tensorflow\python\framework\op_def_library.py:263: colocate_with (from tensorflow.python.framework.ops) is deprecated and will be removed in a future version.
Instructions for updating:
Colocations handled automatically by placer.

WARNING:tensorflow:From C:\Users\thisi\AppData\Local\Continuum\anaconda3\lib\site-packages\keras\backend\tensorflow_backend.py:3445: calling dropout (from tensorflow.python.ops.nn_ops) with keep_prob is deprecated and will be removed in a future version.
Instructions for updating:
Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.

```
In [ ]: ►
```

```
In [43]: ▶ rnn_model.fit(x_train_time_series_data, y_train_time_series_data,
                        batch_size=BATCH_SIZE, epochs=1000, shuffle=False,
                        validation_data=(x_test_time_series_data, y_test_time_series_data), verbose=2)
```

```
Epoch 849/1000
- 23s - loss: 133.4733 - val_loss: 37.4652
Epoch 850/1000
- 23s - loss: 145.5879 - val_loss: 36.4476
Epoch 851/1000
- 23s - loss: 128.7842 - val_loss: 59.1308
Epoch 852/1000
- 22s - loss: 137.5519 - val_loss: 70.2305
Epoch 853/1000
- 23s - loss: 134.0128 - val_loss: 69.4804
Epoch 854/1000
- 23s - loss: 133.6396 - val_loss: 77.5564
Epoch 855/1000
- 23s - loss: 141.4769 - val_loss: 30.5555
Epoch 856/1000
- 22s - loss: 134.5154 - val_loss: 158.0352
Epoch 857/1000
- 22s - loss: 140.4537 - val_loss: 37.4736
Epoch 858/1000
- 22s - loss: 136.2424 - val_loss: 61.5758
```

```
In [44]: ▶ rnn_model.summary()
```

Layer (type)	Output Shape	Param #
gru_1 (GRU)	(None, 60, 3)	81
gru_2 (GRU)	(None, 3)	63
dense_1 (Dense)	(None, 5)	20
dropout_1 (Dropout)	(None, 5)	0
dense_2 (Dense)	(None, 1)	6

```

=====
Total params: 170
Trainable params: 170
Non-trainable params: 0
=====
```

```
In [45]: ▶ rnn_model.save('./models/RNN_GRU_stock_price_prediction_model.h5')
```

```
In [46]: ▶ train_predict= rnn_model.predict(x_train_time_series_data)
```

```
In [47]: ▶ x_train_time_series_data.shape
```

```
Out[47]: (12960, 60, 5)
```

```
In [48]: ▶ train_num_date_corresponding_to_y.shape
```

```
Out[48]: (12960,)
```

```
In [49]: ▶ y_train_time_series_data
```

```
Out[49]: array([ 5.9787,  5.9622,  5.9037, ..., 165.22 , 165.88 , 163.   ])
```

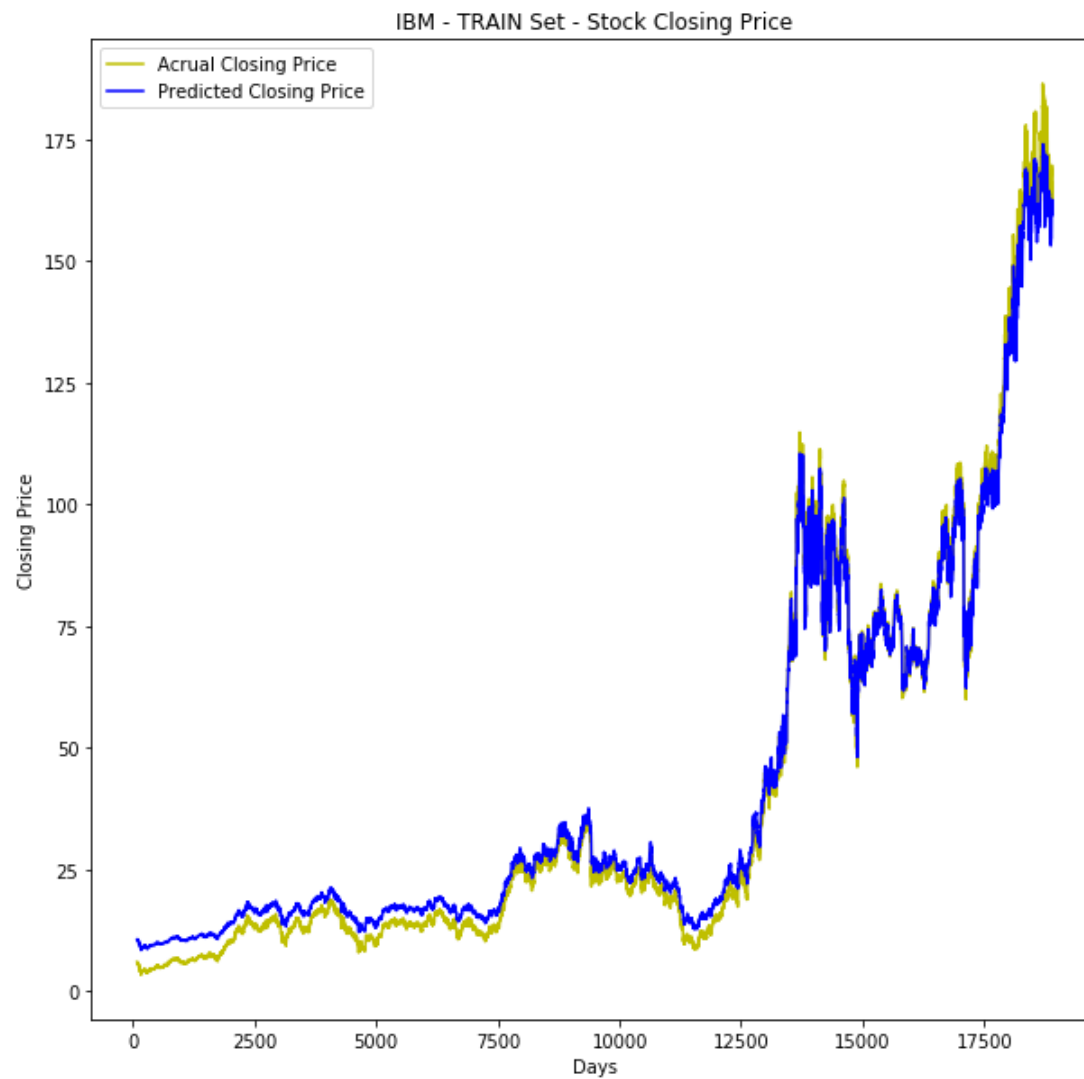
```
In [50]: ▶ y_train_time_series_data
```

```
Out[50]: array([ 5.9787,  5.9622,  5.9037, ..., 165.22 , 165.88 , 163.   ])
```

```
In [64]: ► fig = plt.figure(figsize=(10,10))
ax = fig.add_subplot(111)
ax.set_title('IBM - TRAIN Set - Stock Closing Price')
#ax.scatter(x=data[:,0],y=data[:,1],label='Data')

plt.plot(train_num_date_corresponding_to_y, y_train_time_series_data, color='y', label='Acrua1 Closing Price')
plt.plot(train_num_date_corresponding_to_y, train_predict.ravel(),color='b', label='Predicted Closing Price')

#plt.plot(data[:,0], m*data[:,0] + b,color='red',label='Our Fitting Line')
ax.set_xlabel('Days')
ax.set_ylabel('Closing Price')
ax.legend(loc='best')
plt.show()
```



```
In [62]: test_predict= rnn_model.predict(x_test_time_series_data)
```

```
In [65]: ► fig = plt.figure(figsize=(10,10))
ax = fig.add_subplot(111)
ax.set_title('IBM - TEST Set - Stock Closing Price')
#ax.scatter(x=data[:,0],y=data[:,1],label='Data')

plt.plot(test_num_date_corresponding_to_y, y_test_time_series_data, color='g', label='Acrua1 Closing Price')
plt.plot(test_num_date_corresponding_to_y, test_predict.ravel(),color='orange', label='Predicted Closing Price')

#plt.plot(data[:,0], m*data[:,0] + b,color='red',label='Our Fitting Line')
ax.set_xlabel('Days')
ax.set_ylabel('Closing Price')
ax.legend(loc='best')
plt.show()
```




https://www.dlology.com/blog/how-to-use-return_state-or-return_sequences-in-keras/ (https://www.dlology.com/blog/how-to-use-return_state-or-return_sequences-in-keras/)

In []: ▶

